

Foundation models for medical image analysis have enabled the realization of an age-old aspiration of biomedical AI: to develop a unified, adaptable system that can be applied across organs, modalities, and pathologies. Ma et al.'s "Segment Anything in Medical Images" (MedSAM) is a prominent example of this trend, but the LLM-generated summary presented is far more enthusiastic than the paper itself in terms of technical properties and suitability for clinical deployment. A careful reading of the original article reveals a more cautious and evidence-based story. Rather than a perfect, entirely autonomous universal segmenter, MedSAM is a prompt-based foundation model whose performance depends critically on its training data, prompting strategy, and evaluation context. Mischaracterizing these aspects is not a semantic quibble: the exaggerations in the summary, if taken at face value in a clinical environment, could translate directly into unsafe automation and misplaced trust.

Some of the technical claims in the summary differ markedly from the methods and results reported by Ma et al. First, the summary calls MedSAM “the first universal foundation model for the automatic segmentation of medical images” that enables “fully autonomous segmentation without any human intervention.” In the paper, the authors indeed position MedSAM as a foundation model trained on a large, heterogeneous corpus of medical images, but within the well-established prompt-based segmentation paradigm inherited from Meta’s Segment Anything Model (SAM). MedSAM produces segmentations conditional on prompts, points, boxes, or other spatial information defining what needs to be segmented. It is exactly this promptability that underpins its flexibility across tasks. Presenting MedSAM as a model that “needs no human input” erases this central design feature and misleads the reader about how the system is intended to work.

Second, the summary claims that MedSAM has “perfect precision” and can “instantly recognize new tumors or organs in clinical data,” even when confronted with pathologies not seen during training. The original paper reports performance using standard quantitative measures such as the Dice similarity coefficient across a wide range of benchmark tasks and datasets. While MedSAM often outperforms task-specific baselines and shows promising generalization, the performance is neither uniform nor perfect: scores vary across organs, modalities, and tasks, and the authors explicitly acknowledge residual errors and failure modes. Elevating these nuanced results to

claims of “perfect precision” misrepresents the evidential weight of the experiments and blurs the boundary between strong empirical performance and idealized performance.

Third, the summary states that MedSAM “completely eliminates the need for time-consuming manual annotations by radiologists” and can serve as a “universal replacement for current medical image segmentation systems.” By contrast, Ma et al. consistently frame MedSAM as a means to substantially reduce annotation burden and enable more effective interactive segmentation, transfer learning, and downstream model development. Even if a foundation model provides high-quality initial masks, expert review remains essential to correct remaining errors, curate training sets, adapt to new domains, and monitor performance over time. The authors do not claim that manual annotation loses its place, nor that MedSAM can simply replace all existing systems in all clinical scenarios. The leap from “foundational and broadly applicable” to “universal replacement” is introduced by the LLM, not by the paper.

These distortions are not harmless when transferred into clinical practice. Consider a radiology department under resource limitations where decision-makers see only this LLM summary. They could reasonably conclude that MedSAM can be used as a complete, automatic, modality-agnostic segmentation engine applicable to any tumor or organ with negligible error and no radiologist supervision. Under such assumptions, there would be a powerful incentive to restructure workflows around near-total automation. Yet real-world medical imaging is full of distribution shifts (different scanners, protocols, and populations), artifacts, and rare or atypical presentations, the very scenarios where model performance is most fragile. If clinicians overtrust the system based on exaggerated claims, missed lesions, inaccurate tumor volumes, or mis-segmented organs may go undetected, with direct consequences for diagnosis, staging, and treatment planning.

Equally problematic is the suggestion that manual annotation is no longer necessary. High-quality human labels are the backbone not only of initial training, but of constant validation, domain adaptation, and post-deployment surveillance. If stakeholders internalise the idea that foundation models have made human annotation redundant, they may underfund the processes required to ensure ongoing safety and fairness. For instance, if MedSAM exhibits systematically lower performance in underrepresented demographic groups or in resource-limited settings with older scanners, those disparities will only be identified through careful, human-led evaluation

and error analysis. The summary's dismissal of limitations as "minor" (mainly computational cost and occasional issues with poor quality images) hides more structural issues of bias, calibration, and uncertainty that are highly relevant for safe clinical use.

The broader problem highlighted here is the role of automated LLM summaries in scientific and clinical workflows. Large language models can synthesise complex material quickly, but they are prone to feeding hype, smoothing away caveats, and producing confident but unsupported statements. In the case of MedSAM, the LLM appears to have taken legitimately ambitious phrases from the paper, such as "universal medical image segmentation" and "foundation model," and inflated them into categorical claims about universal replacement, perfect precision, and total autonomy. For a non-expert reader, these embellishments are difficult to spot, yet they shift the perceived contribution from "strong, broadly applicable foundation model with promising generalization" to "problem solved."

In scientific work, normalising this kind of automated overstatement has several risks. Researchers who rely on LLM summaries to scan literature may unintentionally perpetuate exaggerated claims, skewing their understanding of the state of the art and potentially misdirecting follow-up studies. Early-stage concepts may be treated as mature technologies, while open issues and limitations receive too little attention. When such distorted narratives creep into grant proposals, reviews, or policy documents, the cumulative effect can be a misalignment between what models are empirically achieving and what the community collectively believes they are achieving. The epistemic health of any field depends on careful calibration between evidence and claims; LLMs can easily disrupt this balance if used uncritically.

In clinical decision-making, the stakes are even higher. As hospitals experiment with AI assistants to summarise evidence or support procurement decisions, there is a real possibility that summaries like the one in this prompt could influence which tools are trialled, how workflows are designed, or how much human supervision is retained. Yet LLMs lack grounded awareness of risk, have no access to real-time performance data, and bear no accountability for downstream harm. Treating their summaries as authoritative descriptions of complex medical AI systems is fundamentally at odds with principles of reliable, human-centered AI. At best, such summaries

should be preliminary, clearly labeled drafts that prompt experts to engage with the primary literature, not substitutes for it.

This critique does not understate the real contributions of MedSAM or the potential utility of LLMs. MedSAM demonstrates that large, heterogeneous medical imaging data can underpin foundation models with non-trivial cross-task generalisation, providing a powerful basis for interactive tools, rapid prototyping, and downstream fine-tuning. Likewise, LLMs can be invaluable for brainstorming, drafting, and making dense articles more accessible, provided their limitations are understood, and their outputs are checked against sources. The key challenge is integrating both technologies into socio-technical systems that preserve human supervision, promote critical thinking, and maintain robust evaluation practices.

In sum, the LLM's description of MedSAM inflates generalisation into universality, strong performance into "perfect precision," prompt-based segmentation into fully automatic operation, and annotation reduction into annotation abolition. In clinical environments, such misinterpretations could encourage overconfident deployment and underinvestment in expert validation. This case illustrates why automated LLM summaries should be considered starting points for human critical appraisal, not authoritative substitutes, particularly when the subject matter bears directly on patient safety and healthcare decision-making.